

Should You Even Train a Neural Network?

A Forecasting-Based Early Rejection Strategy under Budget Constraints

Santiago Ramírez Castañeda

Universidad de La Sabana

Faculty of Engineering

M.Sc. in Applied Analytics

Thesis Advisor: **Félix Mohr**

October 10, 2025



Abstract

Neural networks have achieved remarkable performance across diverse tasks, yet their design and training remain computationally demanding due to vast architectural search spaces and long convergence times. This thesis investigates whether early-training signals can reliably predict when neural architectures are worth pursuing relative to strong non-neural baselines. Building on this premise, we propose *ShouldYouEvenNN*, a decision-oriented framework that re-frames Neural Architecture Search (NAS) as a probabilistic decision problem rather than a purely optimization task. The framework leverages early learning dynamics and lightweight forecasting to estimate the likelihood that a neural candidate will surpass non-neural models, integrating this evidence into a binary decision mechanism that determines whether further neural training is justified. that weighs potential performance gains against projected computational cost. Through this mechanism, *ShouldYouEvenNN* selectively continues training only when early evidence suggests a positive expected return, effectively governing when neural exploration is justified. Empirical evaluations across 71 OpenML datasets demonstrate that the framework can forecast and select architectures approaching oracle-level performance while remaining competitive with ensemble-based baselines at a fraction of the computational cost. These findings indicate that early-training information, when modeled probabilistically, provides a robust and reproducible foundation for fast, cost-aware neural architecture search.

Chapter 1

Introduction

This chapter outlines the motivation for the study and frames the central question: *when is training a neural network worth the computational cost?* It summarizes the limitations of conventional NAS methods, the dominance of non-neural baselines, and the rationale for reframing architecture search as a decision problem.

Neural networks have transformed modern machine learning, achieving state-of-the-art performance across diverse domains such as vision, language, and forecasting. Yet their design and training remain expensive and limiting. The vastness of possible architectures, hyperparameters, and learning dynamics makes neural exploration a computational gamble—especially when each evaluation demands full convergence to reveal its true potential.

Paradoxically, in structured or tabular domains -the setting of most practical data analytics- non-neural methods still dominate (Shwartz-Ziv & Armon, 2022). Ensemble-based models such as Gradient Boosted Trees and Random Forests routinely outperform neural networks under comparable optimization budgets, offering strong generalization with minimal tuning and far lower computational cost (Borisov et al., 2022; Grinsztajn et al., 2022; Mohr & van Rijn, 2023). This persistent asymmetry raises a fundamental and largely unaddressed question: *when is training a neural network actually worth it?*

Automated Machine Learning (AutoML) frameworks have made great strides toward automation and efficiency. They identify strong baseline models through algorithm selection and hyperparameter optimization, often without requiring deep expertise (Elsken et al., 2019;

Mohr & Wever, 2023). However, while AutoML systems streamline the discovery of high-performing non-neural configurations, Neural Architecture Search (NAS) extends this automation to deep models—frequently at the expense of massive computational budgets. Based on our initial review, most NAS research focuses on accelerating search within the neural family, rather than determining whether neural architectures are even competitive for the problem at hand.

Recent work has shown that early training dynamics—partial learning curves, gradient signals, or batch-level metrics—can reveal predictive information about a model’s eventual performance (Domhan et al., n.d.; Egele et al., 2024; Ren et al., 2022). These findings suggest that evaluating a candidate’s promise does not necessarily require full training. What remains missing is a principled mechanism to *act* on this information: a decision rule that weighs early neural evidence against efficient non-neural baselines.

This thesis introduces *ShouldYouEvenNN*, a forecasting-driven decision framework designed to answer precisely that question. Instead of treating NAS as a blind search for optimal architectures, *ShouldYouEvenNN* formulates it as a rational decision process under computational uncertainty. Each neural candidate is allocated a uniform and limited amount of training effort, from which early validation signals are extrapolated to forecast long-term accuracy. These forecasts are then evaluated probabilistically against a reference baseline through an expected-utility criterion that quantifies whether continued neural training is justified.

By transforming early learning behavior into actionable evidence, *ShouldYouEvenNN* bridges the gap between AutoML efficiency and NAS exploration. The remainder of this work situates *ShouldYouEvenNN* within the broader landscape of automated model design and selection. The next section reviews the evolution of AutoML systems, emphasizing how they streamline pre-processing, algorithm selection, and hyperparameter optimization to identify competitive non-neural baselines efficiently. It then contrasts these approaches with NAS methods, which aim to automate neural design but remain constrained by high computational costs. Finally, it surveys recent advances in early-termination and learning-curve forecasting strategies that seek to reduce such costs—developments that directly motivate the decision-oriented formulation adopted by *ShouldYouEvenNN*.

Chapter 2

Background and Related Work

This chapter situates the research within the existing literature on AutoML, NAS, and early-signal forecasting. It traces the progression from automated model selection to architecture search, emphasizing efficiency and decision-making under uncertainty.

2.1 Automated Machine Learning and Model Selection

AutoML aims to make advanced modeling broadly accessible by automating the stages of data preprocessing, model selection, and hyperparameter optimization. By formalizing these steps as a joint optimization problem, AutoML systems remove much of the trial-and-error traditionally required in machine learning workflows (Baratchi et al., 2024). A central theoretical formulation within AutoML is the *Combined Algorithm Selection and Hyperparameter Optimization* (CASH) problem (Thornton et al., 2013), defined as

$$A^*, \lambda^* \in \arg \min_{A^{(j)} \in \mathcal{A}, \lambda \in \Lambda^{(j)}} \frac{1}{k} \sum_{i=1}^k \mathcal{L} \left(A_{\lambda}^{(j)}, D_{\text{train}}^{(i)}, D_{\text{valid}}^{(i)} \right),$$

where $\mathcal{A}$ denotes the set of learning algorithms, $\Lambda^{(j)}$ their hyperparameter spaces, and $\mathcal{L}$ a validation loss function. This definition formalizes the automated search for the optimal algorithm–configuration pair (A^*, λ^*) that minimizes validation error over k folds.

Early AutoML frameworks approached the CASH problem using Bayesian optimization and ensemble strategies to traverse algorithm and hyperparameter spaces efficiently (Nguyen

et al., 2022; van Hoof & Vanschoren, 2021; Z. Wang et al., 2017; Wu et al., n.d.). Subsequent generations integrated meta-learning and evolutionary search (Jurado & Vanschoren, 2022; Salehin et al., 2024), balancing exploration and exploitation to reduce runtime while maintaining accuracy. Recent work, by Mohr and Wever (2023) demonstrated that much of this complexity can be reduced through separable optimization: optimizing preprocessing and modeling stages independently yields comparable performance to full CASH solvers, revealing that simplicity and transparency often outperform algorithmic intricacy.

AutoML’s evolution reflects a clear principle: efficient automation can replace exhaustive experimentation without sacrificing generalization. This principle provides the foundation upon which neural automation—specifically, NAS—was later built.

2.2 Neural and Non-Neural Baselines on Tabular Data

Despite the success of neural networks in unstructured domains such as vision, text, and speech, their dominance does not extend to tabular or structured data. Comprehensive reviews and benchmarks consistently show that tree-based ensembles, particularly Gradient Boosted Decision Trees (GBDTs) and Random Forests, outperform feedforward neural networks on the majority of tabular datasets (Borisov et al., 2022; Grinsztajn et al., 2022; Mohr & van Rijn, 2023). The reason is both architectural and statistical: GBDTs naturally handle heterogeneous feature scales, categorical variables, and missing values without extensive preprocessing, whereas neural networks require larger datasets, hyperparameter sensitivity, and extensive regularization to achieve comparable results.

Nonetheless, this superiority is not absolute. McElfresh et al. (2023) showed that deep networks can outperform GBDTs in high-dimensional and interaction-heavy datasets, provided sufficient data and regularization. Gardner et al. (2022) found that ensembles often yield better subgroup robustness, while neural networks can capture nonlinear feature hierarchies inaccessible to additive models. Also, Gorishniy et al. (2021) found no universally superior solution when testing GBDTs versus models such as ResNet or Transformers. The tension between flexibility and efficiency persists: neural models can, in principle, surpass non-neural ones,

but only under favorable conditions that justify their cost. This tension motivates the central question underlying this work.

2.3 Neural Architecture Search and the Problem of Efficiency

NAS interprets the AutoML approach now focused for neural networks. For a neural architecture A parameterized by weights θ , the NAS objective seeks to minimize validation loss, but that is subject to be done by the optimization of its parameters:

$$\min_{A \in \mathcal{A}} \mathcal{V}(A, \theta^*(A)) \quad \text{subject to} \quad \theta^*(A) = \arg \min_{\theta} \mathcal{L}(A, \theta),$$

where $\mathcal{V}$ denotes validation performance and $\mathcal{L}$ the training loss. This highlights NAS’s inherent cost: evaluating $\mathcal{V}$ requires full training for every candidate A .

NAS approaches have relied on reinforcement learning and evolutionary algorithms (Mendoza et al., 2016; Zoph & Le, 2016), but computational demands were immense—training tens of thousands of networks from scratch on massive datasets. Later works introduced weight sharing and one-shot models (Jin et al., 2019; Lu et al., 2024), as well as differentiable search spaces (Liu et al., 2024; W. Wang & Li, 2025), drastically reducing redundant computation. Hybrid methods such as Lin et al. (2024), Sakamoto et al. (2023), and Zheng et al. (2021) further improved exploration through multi-armed bandits and task transfer.

Still, scalability and reproducibility remain open issues (Mohr & van Rijn, 2022). Even efficient NAS systems are often constrained by the cost of evaluating candidates to convergence, making them impractical for small or medium-sized problems. Ren et al. (2022) and Garcia-Garcia et al. (2025) proposed progressively expanding training budgets to reduce wasted computation, while Jurado and Vanschoren (2022) introduced multi-fidelity strategies that evaluate partial models before committing additional resources. Despite these advances, most approaches still assume that some form of full training will ultimately be completed. The missing element is a mechanism to determine whether such training should even begin.

2.4 Early-Training Signals and Lightweight Evaluation

Research has shown that early training signals (e.g. recorded after only a few epochs or even mini-batches) can be predictive of final model quality. Learning-curve extrapolation methods such as Domhan et al. (n.d.) and Swersky et al. (2014) fit exponential, logistic, or rational functions to partial validation trajectories, estimating their asymptotic performance. Subsequent work by White et al. (2021), Wistuba and Pedapati (2020), and Yan et al. (2021) expanded this approach to NAS benchmarks, showing that forecast accuracy correlates strongly with final validation outcomes.

Alternative methods exploit static or gradient-based indicators: Wen et al. (2019) introduced zero-cost proxies based on Jacobian conditioning, while Ru et al. (n.d.) used gradient norms to predict overfitting potential. Probabilistic and Bayesian approaches (Adriaensen et al., 2023; Wei et al., 2020) combine these signals into uncertainty-aware forecasts, accelerating NAS by pruning underperforming candidates early. Parallel research in AutoML applied similar ideas to time-constrained model search, using partial evaluations and meta-surrogates to approximate performance across pipelines (Aach et al., 2023; Guan et al., 2025; Solís-Martín et al., 2025).

The idea of halting or downscaling unpromising runs has also been applied to hyperparameter optimization. Egele et al. (2023, 2024) demonstrated that one-epoch performance can approximate the outcome of extensive search, with negligible accuracy loss and substantial time savings. Such findings suggest that early metrics are not merely training diagnostics but informative variables—signals from which future performance can be inferred.

2.5 Forecasting and Decision-Making in NAS

As seen, forecasting has become an integral part of efficiency-driven NAS. Bayesian curve-fitting (Adriaensen et al., 2023), Gaussian-process surrogates, and polynomial extrapolations (Mebirouk et al., 2024; Patel et al., 2024) enable fast ranking of candidates based on partial data. W. Wang and Li (2025) and Nandal et al. (2025) introduced uncertainty quantification into this process, allowing NAS algorithms to balance exploration and exploitation adaptively.

However, even the most advanced forecasters typically serve as auxiliary tools for ranking, not as explicit decision-makers. The question of whether neural exploration is worth continuing at all—remains largely heuristic.

Summary. Together, these sections reveal a persistent gap: existing NAS methods optimize *what* to train but not *whether* to train, motivating the decision-oriented formulation of *ShouldYouEvenNN*.

Chapter 3

Problem Definition and Motivation

This chapter defines the conceptual gap between neural and non-neural automation. It formalizes the decision question—under what conditions neural training becomes rational— and introduces the expected-utility formulation that will guide later chapters.

The literature reviewed above reveals a clear imbalance in how automation has evolved across learning paradigms. While AutoML has successfully optimized model selection and hyperparameters across non-neural families, NAS remains heavily focused on structural optimization rather than strategic decision-making. That is, NAS seeks the best neural architecture *assuming* neural models are worth training, whereas AutoML frameworks often demonstrate that simpler, non-neural models suffice for most real-world problems (Borisov et al., 2022; Grinsztajn et al., 2022; Mohr & van Rijn, 2023). This assumption of neural worthiness produces a systematic inefficiency: NAS frameworks spend vast computational budgets exploring architectures that, for many datasets, cannot outperform ensemble-based baselines.

The challenge, therefore, is not merely to accelerate NAS, but to determine when neural training is rational in the first place. Multi-fidelity and early-termination techniques (Egele et al., 2023, 2024; Garcia-Garcia et al., 2025; Jurado & Vanschoren, 2022; Ren et al., 2022) mitigate cost by pruning unpromising runs; yet these strategies remain reactive—they halt low-performing candidates but do not decide whether neural exploration itself is justified. The field lacks a unified mechanism for comparing early neural progress against established non-neural standards to decide whether additional neural optimization is worthwhile.

This gap motivates a reframing of NAS as a *decision problem under uncertainty*. Rather than searching exhaustively within a neural space, the goal is to use partial training evidence to estimate the probability that neural models can surpass non-neural baselines, and to weigh that probability against the expected computational cost of continuing. This view treats model training not as a guaranteed investment but as a stochastic process whose continuation should depend on its expected return.

Let V_{NN} denote the forecasted validation performance of a neural architecture and V_{base} the performance of a strong non-neural baseline such as a Gradient Boosted ensemble. The central question can then be expressed simply as a binary decision:

Does there exist any parameterization θ such that $V_{NN}(\theta) > V_{base}$?

In essence, the framework replaces the question *What is the optimal architecture?* with a simpler and more actionable one: *Is there a neural candidate worth training further at all?*

Instead of focusing exclusively on architecture search or hyperparameter tuning, *ShouldYouEvenNN* introduces a decision layer that asks, given limited evidence, whether investing additional computation in neural training is expected to pay off. This perspective directly integrates insights from AutoML optimization (Baratchi et al., 2024; Mohr & Wever, 2023; Thornton et al., 2013), NAS efficiency research (Lin et al., 2024; Liu et al., 2024; Mendoza et al., 2016; W. Wang & Li, 2025; Zoph & Le, 2016), and early-signal forecasting studies (Adriaensen et al., 2023; Domhan et al., n.d.; Wei et al., 2020; White et al., 2021).

In this sense, *ShouldYouEvenNN* does not compete with NAS—it governs it. It serves as a probabilistic decision mechanism capable of halting neural exploration when the evidence suggests that the expected payoff no longer justifies the computational cost. This conceptual shift transforms neural architecture search from an open-ended optimization procedure into a disciplined, empirically-driven decision process, aligning model discovery with the principles of rational resource allocation.

3.1 Research Question

The core question behind this research is: *Can early-training signals and lightweight evaluation strategies be effectively used to predict whether neural network architectures will achieve competitive performance compared to strong non-neural models, enabling a fast and resource-efficient NAS framework?*

3.2 Research Objectives

3.2.1 General Objective

To develop a fast and resource-efficient NAS framework capable of determining whether neural networks are likely to yield competitive performance compared to strong non-neural models, by leveraging early training information and lightweight evaluation strategies.

3.2.2 Specific Objectives

- Develop a lightweight decision mechanism that classifies neural candidates as either worth continuing or stopping, based on early learning behavior relative to strong non-neural baselines.
- Develop a NAS framework that determines, through early evidence, whether neural exploration is justified for a given task, prioritizing efficiency and reliability over exhaustive search.
- Empirically validate the framework by benchmarking its capacity to identify tasks where neural networks provide a measurable advantage over non-neural models, evaluating both predictive performance and computational efficiency.

Summary. This framing transforms NAS from a brute-force optimization exercise into a resource-allocation problem, setting the theoretical foundation for the proposed methodology.

Chapter 4

Methodology

4.1 Overview of the ShouldYouEvenNN Framework


ShouldYouEvenNN is a decision-oriented framework that reframes neural architecture search (NAS) as a rational continuation problem under limited computational budgets. Instead of exhaustively optimizing architectures, the framework evaluates whether further neural training is *worth pursuing* based on early training evidence. It integrates three ideas: (i) lightweight learning-curve forecasting, (ii) probabilistic scoring, and (iii) a cost-aware decision rule derived from expected utility.

Figure 4.1 provides a schematic view of the overall process. The system operates in sequential phases:

1. **Data preprocessing.** Input datasets are standardized and split deterministically into training, validation, and test subsets. These serve as the foundation for both neural and non-neural evaluations.
2. **Candidate initialization.** A bounded search space of multilayer perceptrons (MLPs) defines the structural and hyperparameter configurations from which neural candidates are instantiated.
3. **Batch-level training.** Each candidate undergoes short, one-epoch training rounds on progressively larger subsets of data, logging batch-level validation metrics that describe its early learning trajectory.

4. **Forecasting.** Early validation curves are extrapolated using shape-aware morphing of rational, sigmoid, and linear models, yielding a forecasted accuracy and its confidence interval.
5. **Probabilistic scoring.** Forecasts are transformed into a probability of surpassing a strong non-neural baseline using a hybrid Monte Carlo–sigmoid estimator, producing a normalized score per candidate.
6. **Evolutionary refinement.** High-probability candidates are recombined and mutated to form the next generation, while unpromising ones are discarded, ensuring both diversity and efficiency.
7. **Expected-utility decision.** At the end of each generation, the top candidate is evaluated through an expected-utility criterion that weighs potential forecasted gain against remaining computational cost. Training continues only when this expected utility is positive.
8. **Refinement and final decision.** If the utility threshold is met, the corresponding architecture is refined under early-stopping; otherwise, a fallback refinement confirms whether neural networks are uncompetitive for the dataset.

This workflow transforms NAS from an open-ended optimization process into a disciplined decision mechanism that determines whether additional neural training is justified. Each stage contributes measurable evidence—data preprocessing ensures comparability, forecasting quantifies belief, and the utility criterion converts that belief into an actionable decision. The result is a reproducible, cost-aware approach that aligns neural exploration with the thesis objective of identifying when neural networks are rationally worth training.



figures/framework_overview.pdf

Figure 4.1: High-level overview of the ShouldYouEvenNN framework. Early batch-level signals are transformed into forecasts, probabilities, and expected utilities that govern whether neural training continues or terminates.

4.2 Core Components

4.2.1 Data Preprocessing Pipeline

To ensure consistency and reproducibility across all datasets, ShouldYouEvenNN employs a deterministic preprocessing pipeline implemented in Python and PyTorch. The pipeline transforms raw data into standardized tensor representations suitable for both neural and non-neural models.

1. **Input handling.** All datasets are loaded into a unified schema with explicit data types. Categorical features are label-encoded, and numerical features are cast to 32-bit floating point values for compatibility with PyTorch tensors.
2. **Missing values.** Missing entries are imputed according to data type: numerical attributes by mean value, categorical attributes by mode. The imputation parameters are computed exclusively from the training split and reused for validation and testing.
3. **Normalization.** Numerical features are standardized to zero mean and unit variance using statistics derived from the training subset only. This normalization is applied consistently across splits to prevent data leakage.
4. **Rare-category handling.** Categories with frequency below a threshold are mapped to a shared `__other__` token to stabilize model training and ensure consistent encoding across folds.
5. **Splitting protocol.** Each dataset is divided into training, validation, and test subsets with an 80/10/10 proportion using stratified sampling for classification tasks. The random seed for splitting is fixed globally to maintain deterministic partitions across experiments.
6. **Tensor conversion.** After preprocessing, all splits are converted into PyTorch tensors and stored under a reproducible random context, ensuring that identical seeds yield identical batch sequences during training.

The preprocessing pipeline guarantees that subsequent evaluations depend solely on model behavior rather than stochastic data artifacts. By enforcing a consistent transformation across all datasets, ShouldYouEvenNN isolates the methodological contribution—forecasting and decision-making—from dataset-specific variance. This separation allows reliable comparison of neural and non-neural outcomes under strictly controlled input conditions.

4.2.2 Candidate Representation and Search Space

In ShouldYouEvenNN, each neural architecture evaluated during the search is represented as a *candidate*. A candidate encapsulates three elements: the model instance, its architectural

specification, and the record of its training and validation metrics. This structure is formalized through the `Candidate` class, which tracks both static attributes (such as layer configuration and optimizer) and dynamic evidence collected during training (batch effort, validation accuracy, forecasted performance, and probabilistic score).

- **Architecture definition.** Each candidate corresponds to a multilayer perceptron (MLP) whose depth, width, activation functions, and regularization parameters are sampled from a bounded search space.
- **Training state.** During the batch-level loop, the candidate accumulates per-batch effort, cumulative validation metrics, and derived indicators such as forecasted accuracy or expected utility.
- **Reproducibility.** All candidates are instantiated under a fixed random context to guarantee that identical seeds produce identical networks and training behavior across runs.

Search Space Design. The search space is deliberately simple yet expressive, balancing structural diversity with computational feasibility. Table 4.1 summarizes the configuration ranges used throughout the framework.

Table 4.1: Search-space configuration for neural candidates.

Hyperparameter	Range / Options
Hidden layers	2–25
Neurons per layer	3–100
Activation functions	ReLU, LeakyReLU, Sigmoid, Tanh, ELU, GELU
Dropout	0.0–0.5 (step 0.1)
Optimizer	Adam, SGD, RMSprop, AdamW
Learning rate	10^{-4} – 10^{-1} (log-uniform)
Batch size	32–1024 (powers of two)
Weight decay	0 – 10^{-2}
Momentum (SGD only)	0.8–0.99
LR scheduler	Step, Exponential, Cosine, None
Shape pattern	Constant, Pyramid, Inverted, Hourglass, Triangular, Irregular

Shape Pattern. The *shape pattern* parameter determines how neuron counts evolve across the hidden layers, influencing representational capacity and gradient flow:

- **Constant:** uniform width across layers.
- **Pyramid:** progressively narrower layers, forming a funnel topology.
- **Inverted:** increasing width toward deeper layers.
- **Hourglass:** compression followed by expansion, similar to autoencoders.
- **Triangular:** width increases then decreases symmetrically around a midpoint.
- **Irregular:** neuron counts sampled randomly within allowed limits to increase architectural diversity.

This parametrization provides sufficient expressive power to explore a wide variety of functional topologies while keeping the computational footprint predictable. Subsequent components of ShouldYouEvenNN exploit this representation to train, forecast, and evolve architectures under strictly controlled budgets.

4.2.3 Batch-Level Training Protocol

ShouldYouEvenNN trains all neural candidates through a short, incremental process that captures informative early learning signals without requiring full convergence. Each candidate undergoes training at the granularity of individual mini-batches, allowing validation metrics to be recorded with high temporal resolution. This design enables accurate forecasting from partial learning curves and ensures that all architectures are evaluated under comparable computational conditions.

1. **Training rhythm.** Every candidate is trained for a single epoch per generation. After each epoch, the portion of the training data used in the next cycle is expanded by a multiplicative growth factor (default 1.4). This preserves the one-epoch rhythm while gradually increasing the diversity of training examples seen by the model.
2. **Batch-level monitoring.** Within each epoch, validation accuracy is measured after a fixed number of mini-batches. The resulting sequence of validation points forms a partial learning curve that characterizes the candidate’s short-term generalization behavior.

3. **Logging.** For each batch, the framework records the current batch index, training loss, validation accuracy, and wall-clock time. These records are stored in the candidate’s internal log and later used by the forecasting module to model the trajectory of validation performance.
4. **Uniform comparison.** Because all candidates share the same batch-level evaluation frequency and deterministic data order, their learning curves are directly comparable across generations, regardless of network size or hyperparameter configuration.

This fine-grained training protocol replaces the conventional epoch-based evaluation common in NAS with a lightweight, evidence-oriented alternative. Rather than committing to long training runs, ShouldYouEvenNN focuses on the early portions of learning where predictive information about future performance is most concentrated. The resulting batch-level trajectories serve as the foundation for the forecasting and decision mechanisms described in the following sections.

4.2.4 Forecasting Module

The forecasting module is responsible for estimating the long-term validation accuracy of a partially trained candidate using only its early learning trajectory. It transforms short sequences of batch-level validation accuracies into predictive models that infer future performance and associated uncertainty. This estimation forms the evidential basis for all subsequent probability and decision computations in ShouldYouEvenNN.

Input and Objective. Given a series of validation accuracies $y = \{y_1, y_2, \dots, y_t\}$ observed after successive mini-batches $x = \{1, 2, \dots, t\}$, the goal is to forecast the accuracy $\hat{y}_f$ expected at a future training horizon T_f . This horizon is defined as

$$T_f = t_{\text{now}} + N \cdot \max(t_{\text{now}}, t_{\text{dataset}}),$$

where t_{dataset} represents the number of batches in one full pass through the data and N controls the extrapolation depth (default $N = 5$).

Shape-Aware Morphing. Neural learning curves evolve through distinct regimes: rapid growth, plateau transition, and eventual stabilization. Relying on a single functional form (e.g., rational or logistic) often fails to capture this diversity. To address this, ShouldYouEvenNN employs a *shape-aware morphing* forecaster that blends three complementary models:

$$f_{\text{morph}}(x) = w_r f_{\text{rational}}(x) + w_s f_{\text{sigmoid}}(x) + w_l f_{\text{linear}}(x),$$

where $f_{\text{rational}}(x) = \frac{ax}{b+x}$, $f_{\text{sigmoid}}(x) = \frac{L}{1+e^{-k(x-x_0)}}$, and $f_{\text{linear}}(x) = mx + c$. The weights (w_r, w_s, w_l) are derived adaptively from the local slope and curvature of the current validation curve, ensuring smooth transitions between acceleration, saturation, and stability phases.

Uncertainty Quantification. For each forecasted value $\hat{y}_f = f_{\text{morph}}(T_f)$, a two-tailed confidence interval is computed as

$$[\hat{y}_{\text{low}}, \hat{y}_{\text{high}}] = [\hat{y}_f - t_\alpha \sigma_y, \hat{y}_f + t_\alpha \sigma_y],$$

where σ_y is the residual standard deviation and t_α the critical value of the Student's t -distribution for significance level α . This interval captures both model uncertainty and noise observed in the partial trajectory.

Forecast Behavior. When few validation points are available, the forecaster defaults to a conservative rule $\hat{y}_f = y_t + 0.05$, preventing over-optimistic projections from insufficient data. As more batches accumulate, the morphing weights stabilize, and the forecast converges toward the empirical trajectory.

Output. The module returns three quantities for each candidate:

1. the mean forecasted validation accuracy $\hat{y}_f$,
2. the lower and upper confidence bounds $[\hat{y}_{\text{low}}, \hat{y}_{\text{high}}]$, and
3. the local slope and variance estimates used to compute the probability of improvement.

This probabilistic forecast converts transient batch-level fluctuations into a smooth, uncertainty-aware prediction of long-term accuracy. In the broader framework, these outputs serve as inputs to the scoring module, which transforms them into probabilistic evidence about whether the neural candidate is likely to surpass the non-neural baseline.

4.2.5 Decision Criterion Based on Expected Utility

ShouldYouEvenNN applies a cost-aware decision rule derived from the principle of expected utility. Rather than treating U as an optimization target, the framework interprets it as a binary indicator that determines whether additional neural training is justified under the current evidence.

Rationale. At the end of each generation, the best-performing candidate is assessed through a scalar quantity that summarizes both its potential gain and its expected computational expense:

$$U = w(p \cdot \Delta) - \lambda C.$$

Here:

- p is the probability that the candidate will surpass the non-neural baseline, estimated from its forecast and confidence interval;
- $\Delta = \hat{y}_f - g$ denotes the forecasted validation-accuracy margin relative to the baseline performance g ;
- C measures the normalized cost of completing the forecast horizon;
- w is an evidence-weighting factor that increases as training progresses, assigning higher credibility to later forecasts;
- λ scales the cost penalty, regulating the tolerance for computational expenditure.

Decision Policy. The framework applies this value as a threshold rule:

Continue training if $U > 0$, otherwise stop.

This simple inequality encodes the continuation policy: training proceeds only when the expected gain, discounted by uncertainty and cost, remains positive. The policy thus implements a rational go/no-go decision rather than an unconstrained search objective.

Interpretation. Although derived from expected-utility theory, this mechanism functions operationally as a conservative binary criterion. It ensures that each dataset receives neural computation only when the accumulated evidence indicates a likely payoff, while automatically halting exploration when expected returns diminish. By embedding this decision boundary within the evolutionary loop, ShouldYouEvenNN transforms NAS into a disciplined judgment process—balancing performance potential against computational prudence.

4.2.6 Refinement Phase

Once the decision criterion signals that further neural training is justified, ShouldYouEvenNN enters the *refinement phase*. This stage focuses on confirming the candidate’s true potential under controlled extended training while keeping total cost bounded.

Triggering. The refinement phase is activated when the continuation condition $U > 0$ is met for the top-ranked candidate. At that moment, the evolutionary search loop halts, and the remaining global time budget is reallocated to fine-tune that candidate. If no candidate ever reaches positive utility before the budget expires, a single fallback refinement is performed on the most promising forecasted model to verify whether neural networks are genuinely uncompetitive for the given dataset.

Training Procedure. During refinement, the selected candidate resumes training under the same data preprocessing, batch order, and hyperparameters used during the search. An early-stopping mechanism monitors the validation accuracy and terminates training when one of the following conditions is satisfied:

1. the validation accuracy surpasses the non-neural baseline metric;
2. no measurable improvement occurs within a fixed patience window; or

3. the remaining time budget is exhausted.

The best validation checkpoint is automatically restored at the end of training to ensure comparability.

Outcome. The refinement yields a definitive judgment for the dataset:

- If at least one refined candidate surpasses the baseline, the dataset is labeled as *neural worthy*;
- otherwise, the system concludes that neural networks are not competitive under the given computational constraints.

Both successful and unsuccessful refinements are logged with their final validation and test accuracies, allowing later analysis of true positive and false negative outcomes.

Role in the Framework. The refinement phase provides empirical closure to the decision process. It validates whether the probabilistic forecast and expected-utility indicator produced a correct continuation call, grounding the theoretical decision in observed performance. Through this final step, ShouldYouEvenNN transforms early statistical evidence into a verifiable experimental conclusion.

4.2.7 Reproducibility and Logging

Reproducibility is a central design principle of ShouldYouEvenNN. Every experiment, candidate, and decision is recorded deterministically to ensure that identical seeds, configurations, and datasets produce identical results. This guarantees that any claim of neural competitiveness can be independently verified.

Deterministic Execution. All random operations—from architecture sampling to evolutionary crossover—are governed by synchronized random generators across NumPy, PyTorch, and Python’s internal state. The framework initializes a global seed before each experiment, ensuring that the same sequence of random numbers is used at every stage. Consequently, rerunning

the search on the same configuration reproduces the exact same population dynamics, forecasts, and final outcomes.

Structured Logging. Each candidate maintains a comprehensive record of its configuration and performance:

- architectural parameters and initialization seed,
- per-batch training and validation metrics,
- forecasted mean and confidence interval,
- probability of surpassing the baseline,
- expected-utility value and continuation decision,
- refinement statistics and final validation/test accuracy.

These records are serialized into structured `.json` and `.csv` ledgers, enabling direct statistical analysis and visual reconstruction of every run.

Ledger Structure. At the end of each generation, the framework consolidates the state of all candidates into a *ledger*. This ledger captures the full evolution of the population, including forecasts, pruning decisions, and utility scores. Because the ledger is produced deterministically, it functions as both a reproducibility artifact and an audit trail of the decision-making process.

Traceability and Transparency. The combination of deterministic randomness, explicit logging, and structured output ensures that the entire search and decision pipeline can be replayed step by step. This transparency is fundamental to the framework’s scientific integrity: ShouldY-ouEvenNN does not rely on stochastic luck but on reproducible evidence guiding whether neural networks are worth training.

Chapter 5

Empirical Evaluation

5.1 Research Questions

The empirical study addresses three complementary questions, each corresponding to one of the framework’s objectives and evaluated through the experiments described in this chapter.

RQ1: Forecast reliability: Can early batch-level validation signals, modeled through the shape-aware morphing forecaster, reliably approximate the final validation performance of neural networks across diverse datasets?

RQ2: Decision effectiveness and efficiency: Can ShouldYouEvenNN distinguish when continued neural training is warranted based on early evidence?

RQ3: External competitiveness: Are the neural architectures selected by ShouldYouEvenNN comparable in predictive performance to those found by high-fidelity NAS benchmarks such as LCBench, despite the framework’s early-termination strategy and strict budget constraints?

5.2 Experimental Setup

The empirical evaluation of ShouldYouEvenNN follows a standardized setup designed to isolate the behavior of its forecasting, decision, and refinement mechanisms from dataset-specific

variability. All experiments were conducted using deterministic seeds and identical preprocessing, ensuring full reproducibility across runs.

5.2.1 Datasets

A total of seventy-one datasets were retrieved from the *OpenML* repository. These datasets span binary and multiclass classification tasks of varying dimensionality, sample size, and class balance, providing a diverse benchmark for evaluating both neural and non-neural performance. Each dataset was preprocessed through the deterministic pipeline described in Section ??, which includes:

- label encoding for categorical variables and standardization for numerical variables,
- mean or mode imputation for missing values,
- stratified 80/10/10 train–validation–test splits, and
- consistent random seeding to guarantee identical partitions across all experiments.

5.2.2 Baseline Models

To establish reference points for comparison, three non-neural baselines were employed:

1. **NaiveAutoML** Mohr and Wever, 2023: a minimal yet competitive AutoML framework that evaluates classical algorithms such as Gradient Boosting, Random Forests, and Logistic Regression without extensive hyperparameter tuning;
2. **HistGradientBoostingClassifier (HGB)**: an optimized tree-based ensemble implemented in `scikit-learn`, used as a strong single-model baseline; and
3. **Standard MLP**: a feedforward neural network trained conventionally to full convergence, serving as a neural reference for unrestricted training cost.

The highest validation accuracy among the non-neural baselines serves as the *goal metric* g , representing the performance threshold that ShouldYouEvenNN attempts to surpass.

5.2.3 Search Configuration

For each dataset, a population of $N = 25$ neural candidates was sampled from the bounded search space defined in Section ???. Each candidate followed the batch-level training protocol described in Section ??, using the following configuration:

- initial data usage of 10% of the training set per candidate,
- single-epoch training per generation with instance growth factor $1.4\times$,
- mini-batch sizes selected from $\{32, 64, 128, 256, 512\}$, and
- synchronized batch order across candidates for fair comparison.

Each generation performed forecasting, scoring, and pruning before generating a new population through mutation and crossover.

5.2.4 Evaluation Protocol

The framework’s decision process was tested under a fixed global time budget:

$$T_{\text{budget}} = \max(T_{\text{HGB}} \cdot \beta_{\text{budget}}, T_{\text{MLP}}, 60 \text{ s}),$$

where T_{HGB} is the training duration of the HGB baseline, $\beta_{\text{budget}} = 3$ a scaling factor for exploration depth, and 60 seconds ensures a minimum runtime for small datasets.

At each generation:

1. validation accuracy curves were collected at batch level;
2. forecasts and confidence intervals were computed using the shape-aware morphing model;
3. probabilities of surpassing the baseline were estimated;
4. expected-utility values were computed for the best candidate;
5. if $U > 0$, refinement was triggered under early stopping; otherwise, the search continued.

All results, including forecast values, confidence intervals, decision outcomes, and training times, were stored in structured ledgers for later aggregation and analysis.

5.2.5 Performance Metrics

Evaluation focused on three complementary dimensions:

- **Forecast fidelity:** mean absolute error (MAE), bias, and RMSE between forecasted and realized validation accuracies.
- **Decision accuracy:** proportion of correct continuation and rejection calls (true positives and true negatives) relative to actual outcomes.
- **Computational efficiency:** total training time and number of candidates processed compared to conventional NAS and baseline methods.

Together, these metrics provide a comprehensive assessment of whether ShouldYouEvenNN delivers reliable, efficient, and reproducible neural decisions across diverse learning scenarios.

5.3 Results

This section presents the empirical results obtained from the experiments described in the previous section, organized according to the three research questions that guide this study. Each subsection addresses one question, summarizing the evidence, findings, and their interpretation in the context of ShouldYouEvenNN’s objectives.

5.3.1 RQ1 – Forecast Reliability

The first research question evaluates whether early batch-level validation signals can reliably forecast the final validation accuracy of neural networks. To test this, 25 random MLP architectures were trained on each dataset for three full passes through the data. At each batch step, validation accuracies were logged, and the shape-aware morphing model produced forecasts of long-term performance together with confidence intervals. The resulting forecasts were compared against the final accuracies obtained after full early-stopping training.

Results. Across 32 representative datasets, the mean absolute error (MAE) between forecasted and realized validation accuracies averaged **0.1465**, with a root mean squared error

(RMSE) of **0.1920** and a mild negative bias of **-0.0324**. Forecasts were therefore slightly conservative, tending to underestimate true performance—a desirable property for a system that must avoid false optimism. The average confidence interval width was **0.26**, with a coverage of **0.56**, indicating that uncertainty was slightly under-calibrated but consistent across datasets. Pearson and Spearman correlations between forecasted and realized accuracies were modest ($r = 0.20$, $p = 0.21$), suggesting that the forecaster captures performance trends rather than precise rankings.

Interpretation. These results confirm that early learning curves contain sufficient predictive information to guide continuation decisions. Although the absolute correlation is low, the conservative bias ensures that ShouldYouEvenNN errs on the side of caution—terminating unpromising candidates early while still identifying genuinely competitive ones. Hence, the forecasting mechanism fulfills its role as a binary decision aid rather than a precise regressor, successfully addressing **RQ1**.

5.3.2 RQ2 – Decision Effectiveness and Efficiency

The second research question examines whether ShouldYouEvenNN makes accurate and efficient continuation decisions compared to full NAS training. This analysis integrates the forecasting, scoring, and decision layers described in the Methodology, evaluating whether the framework correctly identifies when neural exploration is worthwhile under limited budgets.

Results. Across the 71 OpenML datasets, the framework marked 38 as *neural worthy*, representing approximately 54% of all cases. Among these, 16 datasets surpassed the non-neural validation baseline, and 8 maintained that advantage on the held-out test sets. False negatives—cases where the framework rejected a candidate that later surpassed the baseline—were rare, demonstrating a conservative bias. Figure 6.6 illustrates how the expected-utility signal varies with feature dimensionality: datasets with moderate dimensionality tended to yield positive decisions, whereas high-dimensional datasets remained below the threshold, indicating diminishing expected returns.

In terms of efficiency, the total runtime per dataset remained close to the baseline training time of a single MLP, achieving up to two orders of magnitude reduction compared to traditional NAS approaches that require full convergence for every candidate.

Interpretation. The results demonstrate that ShouldYouEvenNN effectively identifies when continued neural training is justified. Its binary decision rule, grounded in conservative probability estimates, prevents wasted computation without missing many genuine opportunities for neural advantage. The trade-off is deliberate: the framework favors reliability and reproducibility over aggressive exploration. Consequently, **RQ2** is supported—ShouldYouEvenNN achieves decision accuracy and computational efficiency consistent with its design goals.




Figure 5.1: Expected-utility signal across datasets with increasing feature dimensionality. Orange markers denote datasets deemed neural worthy ($U > 0$); blue markers represent rejected cases. The clear separation around the $U = 0$ boundary reflects the framework’s conservative decision behavior.

5.3.3 RQ3 – External Competitiveness

The third research question evaluates whether the neural architectures identified by ShouldYouEvenNN remain competitive with high-fidelity NAS results obtained from the Learning Curve Benchmark (LCBench). This comparison assesses the external validity of the framework’s lightweight approach.

Results. Across 35 overlapping datasets between OpenML and LCBench, the best ShouldYouEvenNN models achieved an average validation accuracy of **0.771**, compared to **0.819** for LCBench—an average deviation of approximately **-0.05**. In 37% of cases (13/35), the ShouldYouEvenNN models matched or exceeded the LCBench-reported accuracy. The remaining differences reflect the framework’s intentionally conservative early-termination policy. Despite operating with a fraction of the computational cost, ShouldYouEvenNN consistently reached accuracies within 5–6 percentage points of the fully trained NAS models.

Interpretation. The proximity of these results confirms that ShouldYouEvenNN captures much of the neural performance frontier attainable by exhaustive NAS, while maintaining high efficiency and full reproducibility. It demonstrates that early, uncertainty-aware decisions can lead to models that are competitive with state-of-the-art NAS outcomes, thereby addressing **RQ3**.

Summary. Collectively, the results across RQ1–RQ3 show that:

- early forecasting provides sufficiently reliable evidence for continuation decisions;
- the decision rule accurately balances benefit and cost, preventing unnecessary computation; and
- the selected neural architectures remain competitive with large-scale NAS baselines despite strict budget limits.

These findings confirm that ShouldYouEvenNN meets its intended objectives of reliability, efficiency, and competitive performance under constrained resources.

Chapter 6

Results

This section reports the empirical findings from the forecasting tests and full SHOULDY-OUEVENNN executions. It analyzes forecast accuracy, decision outcomes, and cross-dataset trends.

6.1 Early Signals Forecasting Strategy

To assess the reliability between the forecasting strategy and the "true" results for each model, the next results were analyzed.

In the Figure 1 and 3, two examples of the learning curves of some of the models trained according to the ShouldYouEvenNN strategy is observed, and that were later tested for two specific datasets (*ada* and *phoneme*). Figure 2 and 4 show the error distribution with its respective MAE.

Table 6.1: Summary statistics of the EBE-NAS results.

	mean	std
MAE	0.1465	0.0700
Bias	-0.0324	0.0948
RMSE	0.1920	0.0860
CI coverage	0.5557	0.3201
Avg CI width	0.2574	0.1271
Pearson r	0.2035	0.3270
Spearman rho	0.2149	0.2863





Figure 6.1: Example of learning curves for several models trained under the ShouldYouEvenNN strategy with their respective forecasted accuracy on the ada dataset. The trajectories illustrate the early-stage validation behavior used for forecasting.

The reliability of the forecasting mechanism was assessed across 32 datasets, summarizing its deviation and calibration with respect to the actual validation accuracies observed after full early-stopping training. As shown in Table 6.1, the mean absolute error (MAE) averaged 0.1465, indicating that forecasts deviated by approximately fifteen percentage points from the true final performance on average. The corresponding root mean squared error (RMSE) of 0.1920 reflects the presence of occasional outliers in which the fitted curves diverged from the eventual validation trajectory, typically associated with noisy or stagnant learning curves.

A mild negative bias of -0.0324 suggests that forecasts tended to underestimate final performance as seen as well in Figure 2, a conservative behavior aligned with the framework’s goal of avoiding unnecessary neural investment. The average confidence interval (CI) coverage was 0.56, below the nominal 95% expectation, revealing that the uncertainty estimates were somewhat overconfident. This is consistent with the observed mean CI width of 0.26, which proved too narrow relative to the empirical forecast error. Nevertheless, this conservative calibration



(a) Error distribution of forecasted validation accuracies with corresponding mean absolute error (MAE).

(b) Scatter plot comparing forecasted and observed validation accuracies.

Figure 6.2: Forecasting performance visualization on the ada dataset.

minimizes false positives during the early rejection phase.

The correlation coefficients between forecasted and actual accuracies were modest ($r = 0.20$, $\rho = 0.21$), implying that while the forecasts captured the general direction of performance evolution, they were not strongly predictive of candidate rankings. This behavior highlights that the forecasting subsystem functions effectively as a binary decision aid—indicating whether further training is justified—rather than as a precise performance predictor.

6.2 ShouldYouEvenNN Performance

The evolution observed in Figure 6.5 exemplifies the internal behavior of the framework during the search phase. The mean validation accuracy progressively increases across generations, accompanied by closely aligned forecasted accuracies, indicating that the forecasting module provides a realistic and conservative estimate of performance trends. The probability of surpassing the baseline (p) rises steadily during the first half of the process, reflecting growing evidence that the population contains promising architectures. However, the expected utility (U) remains near zero or slightly negative, showing that the projected computational cost continues to outweigh the marginal expected gain—hence no refinement is triggered at this stage.




Figure 6.3: Example of learning curves for several models with their respective forecasted accuracy on the phoneme dataset.


The cumulative compute curve increases almost linearly, confirming that the training effort per generation is stable and well-bounded by the global budget. Overall, this plot illustrates a typical non-triggering scenario where learning signals improve and forecasts remain well calibrated, yet the expected-utility criterion correctly prevents unnecessary extended training.

After validating the forecasting module, the full algorithm was executed across all benchmark datasets. The aggregated outcomes are summarized in Table 6.2.

Metric	Value
Total Datasets	71
Neural Worthy	38
Validation Surpassed	16
Test Surpassed	8

Table 6.2: Overall results of the ShouldYouEvenNN evaluation across all datasets.

Across the 71 OpenML datasets evaluated, the ShouldYouEvenNN framework identified neural networks as *worth training* in 38 cases, representing just over half of the total sample.



(a) Error distribution of forecasted validation accuracies with corresponding mean absolute error (MAE).

(b) Scatter plot comparing forecasted and observed validation accuracies.

Figure 6.4: Forecasting performance visualization on the phoneme dataset.

Among these, 16 datasets achieved validation accuracies that surpassed their respective non-neural baselines, and 8 maintained this advantage on the held-out test sets.

These results indicate that while *ShouldYouEvenNN*'s decision rule is selective, it remains capable of identifying scenarios in which neural networks provide tangible benefits over strong ensemble-based baselines. The gap between the number of neural-worthy cases and those that ultimately surpassed validation and test metrics illustrates the framework's conservative stance: only a subset of early promising signals translated into verifiable generalization gains. This outcome aligns with *ShouldYouEvenNN*'s primary objective—to allocate additional training effort only when statistically justified by early evidence—thereby reducing wasted computation without systematically overlooking genuinely competitive neural configurations.

Figure 6.6 illustrates how the expected utility assigned by the decision rule varies with dataset feature dimensionality (in logarithmic scale). Each point represents a dataset, where orange markers correspond to cases deemed *neural worthy* (decision = True) and blue markers to those rejected. Cross-shaped markers further indicate instances in which the validation baseline was actually surpassed after extended training.

The plot reveals a consistent separation between positive and negative utility regions. Most




Figure 6.5: Generation-level dynamics for Dataset 3. The figure summarizes how validation accuracy, forecasted performance, probability of surpassing the baseline, and expected utility evolved through successive generations of the *ShouldYouEvenNN* framework.

datasets with positive expected utility ($U > 0$) cluster below a few hundred features, suggesting that *ShouldYouEvenNN* tends to favor neural architectures when the feature space is relatively low-dimensional and structured. Conversely, for high-dimensional datasets (beyond 10^3 features), expected utility values remain largely negative, implying that the computational cost or instability outweighs potential gains in these scenarios.

Interestingly, several datasets near the decision boundary ($U \approx 0$) correspond to borderline cases where the neural models approached but did not exceed baseline performance. This reinforces the probabilistic nature of the decision function: small variations in early learning dynamics can tip the forecast from rejection to acceptance. The figure thus encapsulates the framework’s core behavior—balancing optimism and cost sensitivity—while highlighting that the conditions under which neural networks become worthwhile remain concentrated in medium-sized, moderately complex datasets rather than in very high-dimensional feature spaces. Also, the cases where the decision was defined to be true and $U < 0$ correspond to sce-




Figure 6.6: Dataset Feature Size vs U

narios where the framework initially failed to identify a task as neural worthy, but after extending the best candidate during the refinement phase, the model ultimately surpassed the baseline metric. The percentage of these cases is very low, which shows the robustness and conservative reliability of the decision mechanism. In other words, false negatives—datasets where *ShouldYouEvenNN* prematurely rejected neural exploration despite eventual success—are rare, but are able to be identified by the tool with this last failguard. This confirms that the expected-utility criterion favors caution, prioritizing the prevention of unnecessary training over speculative optimism.

6.3 External Validation against LCBench

To assess whether the architectures identified by *ShouldYouEvenNN* approach the quality of fully trained neural networks reported in standardized NAS benchmarks, we compared their final validation accuracies to those of the **Learning Curve Benchmark (LCBench)** (Zimmer et al., 2020). LCBench includes performance records for hundreds of neural configurations per dataset, providing an empirical ceiling for expected neural performance.

Table 6.3 summarizes the aggregated statistics of the comparison across 35 overlapping datasets.

Table 6.3: Comparison of final validation accuracies between *ShouldYouEvenNN* and LCBench.

Metric	Mean	Std	Min	25%	50%	75%	Max
EBE Final Val. Acc.	0.771	0.231	0.006	0.712	0.819	0.926	0.994
LCBench Final Val. Acc.	0.819	0.155	0.290	0.743	0.838	0.945	0.997
Difference (EBE – LCBench)	-0.048	0.159	-0.871	-0.037	-0.004	0.008	0.097

Across datasets, the average deviation between the best *ShouldYouEvenNN* model and its LCBench counterpart was approximately -0.05 , with standard deviation 0.16 . This indicates that the proposed framework, despite its strict time and data constraints, achieves validation accuracies within five percentage points of the high-fidelity LCBench results on average.

A contingency analysis of surpassing outcomes is shown in Table 6.4.

Table 6.4: Cross-tabulation of datasets where *ShouldYouEvenNN* and LCBench surpass their respective baselines.

LCBench surpassed	<i>ShouldYouEvenNN</i> surpassed baseline	Count
False	False	21
True	False	8
True	True	5
False	True	1

Overall, in roughly 37% of the overlapping datasets (13/35), the final models obtained by *ShouldYouEvenNN* matched or exceeded the validation accuracy reported by LCBench. The remaining discrepancies are modest and primarily reflect the framework’s intentionally conservative stopping policy. These results validate that the early-forecast-driven selection mechanism can recover high-performing neural architectures that remain competitive with exhaustive NAS baselines while preserving computational efficiency.

Summary. The results confirm that early learning dynamics carry sufficient information to guide resource-efficient NAS decisions and that the framework’s conservative bias effectively prevents wasted training.

Chapter 7

Discussion

This chapter interprets the experimental results in light of the research objectives. It discusses reliability, conservativeness, computational efficiency, and alignment with current NAS literature.

7.1 Interpretation and Theoretical Positioning

The results of *ShouldYouEvenNN* validate its central hypothesis: that early training dynamics contain sufficient predictive information to make rational, cost-aware decisions about whether neural networks are worth pursuing. Across the OpenML benchmark, the framework consistently identified competitive architectures without relying on exhaustive convergence, confirming that lightweight forecasting and selective continuation can substitute for full training in many cases.

These findings directly address the long-standing imbalance between AutoML and NAS efficiency. As noted by Baratchi et al. (2024) and Mohr and Wever (2023), AutoML systems achieve scalable performance through separable and interpretable optimization, whereas classical NAS pipelines—often based on reinforcement or evolutionary search (Mendoza et al., 2016; Zoph & Le, 2016)—still assume that all candidates warrant full convergence. *ShouldYouEvenNN* challenges this assumption by introducing an explicit decision layer that governs when neural exploration itself is justified, aligning with the resource-rational principles discussed by Mohr and van Rijn (2022). This integration of decision theory with early-signal

evidence redefines NAS as a process of *judgment under uncertainty* rather than blind optimization.

7.2 Forecasting Reliability and Early-Signal Evidence

The batch-level forecasting component proved both robust and informative. Its morphing curve model, which blends rational, sigmoid, and linear regressors, captured meaningful learning trends even under noisy, low-data conditions. This empirically supports prior claims that partial learning curves are predictive of final outcomes (Domhan et al., n.d.; White et al., 2021; Wistuba & Pedapati, 2020). However, unlike those studies—where forecasts serve as passive ranking tools—*ShouldYouEvenNN* operationalizes them as actionable evidence within a decision pipeline.

The conservative bias observed in the fidelity benchmark is consistent with the cautionary calibration suggested by Egele et al. (2024), who showed that one-epoch signals can be effective predictors when optimism is moderated. By penalizing steep or volatile learning curves, the forecasting mechanism minimizes false positives, favoring stability and reproducibility. This confirms that early-stage training signals can function as probabilistic priors guiding resource allocation, not merely as performance predictors.

7.3 Decision Criteria and Conservativeness

ShouldYouEvenNN evaluates neural candidates through a binary decision rule rather than a formal probabilistic optimization. The goal is not to compute an exact utility or posterior probability, but to decide, as early as possible, whether a given neural configuration is worth further training compared to a strong non-neural baseline.

The framework treats this as a classification-like prediction task: given early evidence—partial learning curves, forecasted validation performance, and approximate training cost—it predicts the outcome of the event ($V_{NN} > V_{base}$). A continuous score S aggregates three normalized components: confidence from the forecast, the estimated improvement margin, and the remaining cost. High or positive values of S indicate that continued neural training is likely

to be beneficial; low or negative values indicate that it is not. This scalar is therefore a surrogate for a yes/no decision rather than a solved objective.

Because the rule is intentionally conservative, *ShouldYouEvenNN* only continues training when there is clear, stable evidence of potential improvement. The decision threshold favors reliability over optimism, ensuring that scarce computation is allocated to candidates with sustained progress instead of transient spikes in validation accuracy. In practice, this conservatism produces a small number of “neural-worthy” datasets, but these cases almost always correspond to verifiable performance gains beyond the baseline. The few false negatives observed—datasets initially rejected but later surpassing the baseline during the fail-safe refinement phase—remain rare and confirm that the rule errs on the side of caution.

This design choice reflects the intended role of *ShouldYouEvenNN*: not as an optimizer that searches for the global best network, but as a filtering mechanism that determines when neural exploration itself is justified. The framework thus prioritizes efficiency and reproducibility over aggressive discovery, embodying a rational yet lightweight approach to early decision-making in neural architecture search.

7.4 Comparison with Baselines and Literature

The selective activation of the “neural worthy” decision primarily on medium-sized and moderately complex datasets mirrors the asymmetry reported by Shwartz-Ziv and Armon (2022), Borisov et al. (2022), and Grinsztajn et al. (2022). Deep models rarely dominate in tabular domains unless nonlinear feature interactions are strong or dataset scale is large, as also observed by McElfresh et al. (2023). *ShouldYouEvenNN* quantifies this boundary empirically, identifying the point at which neural exploration becomes rational given cost constraints.

External comparison with the LCBench benchmark (Zimmer et al., 2020) further positions the framework within the NAS landscape. Despite its lightweight evaluation regime, *ShouldYouEvenNN* achieved validation accuracies within approximately five percentage points of LCBench results on average, and matched or exceeded them in over one-third of shared datasets. This proximity demonstrates that early-evidence forecasting can recover high-quality

architectures typically found through exhaustive search, supporting the general efficiency conclusions of Egele et al. (2023) and Wei et al. (2020).

7.5 Practical Implications and Future Work

Beyond its methodological contributions, *ShouldYouEvenNN* offers practical value for automated pipelines that must decide whether neural training is justified at all. By turning early evidence into a decision signal, it reduces energy consumption and computational waste without systematically discarding competitive candidates. Moreover, the analysis of datasets that consistently trigger positive expected utility could inform meta-learning models capable of predicting ‘neural worthiness’ a priori.

Future research should explore adaptive calibration of the expected-utility parameters (e.g., cost scale and evidence ramp) based on dataset meta-features, and extend the framework to other architectures such as CNNs or Transformers. Integrating such decision layers into end-to-end AutoML systems would advance the field toward a fully autonomous modeling paradigm—one that not only trains models, but also determines, with statistical rigor, when training is worth pursuing.

Summary. Taken together, the discussion supports the central thesis: early-signal forecasting combined with cost-aware utility provides a rational basis for deciding when neural exploration is worthwhile.

Chapter 8

Conclusion

The final chapter synthesizes the conceptual and empirical contributions of the work, highlighting how SHOULDYOU EVENNN reframes NAS as a probabilistic decision process.

The *ShouldYouEvenNN* framework demonstrates that early-training information, when properly modeled and integrated into a probabilistic decision process, can substantially reduce the computational burden of neural architecture search without compromising predictive reliability. By treating the continuation of neural training as a decision under uncertainty rather than a fixed commitment, the framework introduces a principled mechanism for determining when neural exploration is justified relative to strong non-neural baselines.

This study redefines the purpose of NAS. Rather than framing architecture discovery as a competition for the most accurate network, *ShouldYouEvenNN* interprets it as a problem of rational evidence allocation—where limited computational resources are invested only when early learning signals indicate a positive expected return. In this sense, NAS becomes not merely an optimization process, but an adaptive inference system guided by the probabilistic value of continuing to learn.

Empirical results confirm that this perspective is both practical and generalizable. The framework consistently identified when neural networks were likely to be competitive, producing selections that closely approached high performing models while maintaining cost efficiency. Forecasting bias remained intentionally conservative, ensuring that decisions favored reliability over speculative optimism. These findings highlight that early-training signals can

serve not only as predictors of performance but as the statistical foundation for rational, data-driven control over the learning process itself.

Beyond its immediate results, this work contributes conceptually to the broader field of automated model design. It bridges the efficiency focus of AutoML with the representational depth of NAS, demonstrating that automation can evolve from search to judgment—from exhaustive exploration toward selective, evidence-based continuation. By unifying forecasting, probabilistic reasoning, and evolutionary refinement within a single decision framework, *ShouldYouEvenNN* establishes a reproducible path toward fast, transparent, and resource-aware neural model discovery.

Future extensions of this research could explore adaptive calibration of the expected-utility parameters, hierarchical decision layers across modalities (e.g., vision or time series), and tighter integration with pipeline-level AutoML systems. Such developments would move the field closer to a fully autonomous modeling paradigm—one that not only automates training, but also decides, with statistical rigor, when training is worth pursuing at all.

Summary. Ultimately, the study demonstrates that automation in deep learning can evolve from exhaustive search toward selective, evidence-based continuation, an approach that values judgment as much as optimization.

Bibliography

- Aach, M., Inanc, E., Sarma, R., Riedel, M., & Lintermann, A. (2023). Optimal Resource Allocation for Early Stopping-based Neural Architecture Search Methods.
- Adriaensen, S., Rakotoarison, H., Müller, S., & Hutter, F. (2023). Efficient Bayesian Learning Curve Extrapolation using Prior-Data Fitted Networks. <https://arxiv.org/pdf/2310.20447>
- Baratchi, M., Wang, C., Limmer, S., van Rijn, J. N., Hoos, H., Bäck, T., & Olhofer, M. (2024). Automated machine learning: past, present and future. *Artificial Intelligence Review*, 57(5), 122. <https://doi.org/10.1007/s10462-024-10726-1>
- Borisov, V., Leemann, T., Seßler, K., Haug, J., Pawelczyk, M., & Kasneci, G. (2022). Deep Neural Networks and Tabular Data: A Survey. *IEEE Transactions on Neural Networks and Learning Systems*, 35(6), 7499–7519. <https://doi.org/10.1109/TNNLS.2022.3229161>
- Domhan, T., Springenberg, T., & Hutter, F. (n.d.). Speeding up Automatic Hyperparameter Optimization of Deep Neural Networks by Extrapolation of Learning Curves.
- Egele, R., Guyon, I., Sun, Y., & Balaprakash, P. (2023). Is One Epoch All You Need For Multi-Fidelity Hyperparameter Optimization? *ESANN 2023 proceedings*, 555–560. <https://doi.org/10.14428/esann/2023.ES2023-84>
- Egele, R., Mohr, F., Viering, T., & Balaprakash, P. (2024). The unreasonable effectiveness of early discarding after one epoch in neural network hyperparameter optimization. *Neurocomputing*, 597, 127964. <https://doi.org/10.1016/J.NEUCOM.2024.127964>
- Elsken, T., Metzen, J. H., & Hutter, F. (2019). *Neural Architecture Search: A Survey* (tech. rep.). <http://jmlr.org/papers/v20/18-598.html>.

- Garcia-Garcia, C., Derbel, B., Morales-Reyes, A., & Escalante, H. (2025). *Speeding up the Multi-objective NAS Through Incremental Learning* (Vol. 15247 LNAI). https://doi.org/10.1007/978-3-031-75543-9{_}1
- Gardner, J., Popović, Z., Schmidt, L., & Allen, P. G. (2022). Subgroup Robustness Grows On Trees: An Empirical Baseline Investigation. <https://arxiv.org/pdf/2211.12703>
- Gorishniy, Y., Rubachev, I., Khrulkov, V., & Babenko, A. (2021). Revisiting Deep Learning Models for Tabular Data. *Advances in Neural Information Processing Systems*, 23, 18932–18943. <https://arxiv.org/pdf/2106.11959>
- Grinsztajn, L., Oyallon, E., & Varoquaux, G. (2022). Why do tree-based models still outperform deep learning on tabular data? <https://arxiv.org/pdf/2207.08815>
- Guan, J., Zhang, F., Liu, J., Du, X., Shen, X., & Shen, X. A. (2025). A Systematic Study on Early Stopping Metrics in HPO and the Implications of Uncertainty. 18(6), 2150–8097. <https://doi.org/10.14778/3725688.3725689>
- Jin, H., Song, Q., & Hu, X. (2019). Auto-Keras: An Efficient Neural Architecture Search System. *Proceedings of the 25th ACM SIGKDD International Conference on Knowledge Discovery & Data Mining*, 1946–1956. <https://doi.org/10.1145/3292500.3330648>
- Jurado, I. C., & Vanschoren, J. (2022). Multi-fidelity optimization method with asynchronous generalized island model for AutoML. *GECCO 2022 Companion - Proceedings of the 2022 Genetic and Evolutionary Computation Conference*, 220–223. <https://doi.org/10.1145/3520304.3528917>
- Lin, Y., Endo, Y., Lee, J., & Kamijo, S. (2024). Bandit-NAS: Bandit sampling and training method for Neural Architecture Search. *Neurocomputing*, 597, 127684. <https://doi.org/10.1016/J.NEUCOM.2024.127684>
- Liu, X., Tian, J., Duan, P., Yu, Q., Wang, G., & Wang, Y. (2024). GrMoNAS: A granularity-based multi-objective NAS framework for efficient medical diagnosis. *Computers in Biology and Medicine*, 171, 108118. <https://doi.org/10.1016/J.COMPBIOMED.2024.108118>

- Lu, S., Hu, Y., Yang, L., Mei, J., Sun, Z., Tan, J., & Song, C. (2024). PHD-NAS: Preserving helpful data to promote Neural Architecture Search. *Neurocomputing*, 587. <https://doi.org/10.1016/j.neucom.2024.127646>
- McElfresh, D., Khandagale, S., Valverde, J., Prasad, V. C., Ramakrishnan, G., Goldblum, M., & White, C. (2023). When Do Neural Nets Outperform Boosted Trees on Tabular Data? *Advances in Neural Information Processing Systems*, 36. <https://arxiv.org/pdf/2305.02997>
- Mebirouk, N., Amrane, M., Messast, S., & Mazouzi, S. (2024). Enhanced analysis of landslide susceptibility mapping in the proximity of main roads in the province of Skikda, Algeria: using NAS for efficient performance and faster processing. *Modeling Earth Systems and Environment*, 10(5), 6449–6474. <https://doi.org/10.1007/s40808-024-02129-6>
- Mendoza, H., Klein, A., Feurer, M., Springenberg, J. T., & Hutter, F. (2016, November). Towards Automatically-Tuned Neural Networks. In F. Hutter, L. Kotthoff, & J. Vanschoren (Eds.), *Proceedings of the workshop on automatic machine learning* (pp. 58–65, Vol. 64). PMLR. https://proceedings.mlr.press/v64/mendoza_towards_2016.html
- Mohr, F., & van Rijn, J. N. (2022). Learning Curves for Decision Making in Supervised Machine Learning – A Survey. <http://arxiv.org/abs/2201.12150>
- Mohr, F., & van Rijn, J. N. (2023). Fast and Informative Model Selection Using Learning Curve Cross-Validation. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 45(8), 9669–9680. <https://doi.org/10.1109/TPAMI.2023.3251957>
- Mohr, F., & Wever, M. (2023). Naive automated machine learning. *Machine Learning*, 112(4), 1131–1170. <https://doi.org/10.1007/s10994-022-06200-0>
- Nandal, P., Mann, P., Bohra, N., Aldehim, G., Elnour, A. A. H., & Allafi, R. (2025). Tropical cyclone intensity estimation based on YOLO-NAS using satellite images in real time. *Alexandria Engineering Journal*, 113, 227–241. <https://doi.org/10.1016/j.aej.2024.10.072>
- Nguyen, D. A., Kononova, A. V., Menzel, S., Sendhoff, B., & Back, T. (2022). An Efficient Contesting Procedure for AutoML Optimization. *IEEE Access*, 10, 75754–75771. <https://doi.org/10.1109/ACCESS.2022.3192036>

- Patel, S., Patel, D., Patel, A., Nanavati, J., Patel, U., Faldu, A., & Shingala, A. (2024). Accurate and Efficient Real-Time Crop and Weed Identification Using YOLO-NAS: A Neural Architecture Search-Optimized Object Detection Model. *International Journal of Engineering Trends and Technology*, 72(10), 35–45. <https://doi.org/10.14445/22315381/IJETT-V72I10P104>
- Ren, P., Xiao, Y., Chang, X., Huang, P.-y., Li, Z., Chen, X., & Wang, X. (2022). A Comprehensive Survey of Neural Architecture Search. *ACM Computing Surveys*, 54(4), 1–34. <https://doi.org/10.1145/3447582>
- Ru, B., Lyle, C., Schut, L., Fil, M., Van Der Wilk, M., & Gal, Y. (n.d.). Speedy Performance Estimation for Neural Architecture Search. <https://github.com/rubinxin/TSE>.
- Sakamoto, K., Ishibashi, H., Sato, R., Shirakawa, S., Akimoto, Y., & Hino, H. (2023). ATNAS: Automatic Termination for Neural Architecture Search. *Neural Networks*, 166, 446–458. <https://doi.org/10.1016/J.NEUNET.2023.07.011>
- Salehin, I., Islam, M. S., Saha, P., Noman, S., Tuni, A., Hasan, M. M., & Baten, M. A. (2024). AutoML: A systematic review on automated machine learning with neural architecture search. *Journal of Information and Intelligence*, 2(1), 52–81. <https://doi.org/10.1016/j.jiixd.2023.10.002>
- Shwartz-Ziv, R., & Armon, A. (2022). Tabular data: Deep learning is not all you need. *Information Fusion*, 81, 84–90. <https://doi.org/10.1016/J.INFFUS.2021.11.011>
- Solís-Martín, D., Galán-Páez, J., & Borrego-Díaz, J. (2025). A Model for Learning-Curve Estimation in Efficient Neural Architecture Search and Its Application in Predictive Health Maintenance. *Mathematics 2025, Vol. 13, Page 555*, 13(4), 555. <https://doi.org/10.3390/MATH13040555>
- Swersky, K., Snoek, J., & Adams, R. P. (2014). Freeze-Thaw Bayesian Optimization. <https://arxiv.org/pdf/1406.3896>
- Thornton, C., Hutter, F., Hoos, H. H., & Leyton-Brown, K. (2013). Auto-WEKA: Combined Selection and Hyperparameter Optimization of Classification Algorithms. <https://arxiv.org/abs/1208.3719>

- van Hoof, J., & Vanschoren, J. (2021). Hyperboost: Hyperparameter Optimization by Gradient Boosting surrogate models. <https://arxiv.org/pdf/2101.02289>
- Wang, W., & Li, H. (2025). NAS FD Lung: A novel lung assist diagnostic system based on neural architecture search. *Biomedical Signal Processing and Control*, 100. <https://doi.org/10.1016/j.bspc.2024.107022>
- Wang, Z., Gehring, C., Kohli, P., & Jegelka, S. (2017). Ensemble Bayesian Optimization. *International Conference on Artificial Intelligence and Statistics, AISTATS 2018*, 745–754. <https://arxiv.org/pdf/1706.01445>
- Wei, C., Niu, C., Tang, Y., Wang, Y., Hu, H., & Liang, J. (2020). NPENAS: Neural Predictor Guided Evolution for Neural Architecture Search. *IEEE Transactions on Neural Networks and Learning Systems*, 34(11), 8441–8455. <https://doi.org/10.1109/TNNLS.2022.3151160>
- Wen, W., Liu, H., Chen, Y., Li, H., Bender, G., & Kindermans, P. J. (2019). Neural Predictor for Neural Architecture Search. *Lecture Notes in Computer Science (including subseries Lecture Notes in Artificial Intelligence and Lecture Notes in Bioinformatics)*, 12374 LNCS, 660–676. https://doi.org/10.1007/978-3-030-58526-6{_}39
- White, C., Zela, A., Ru, B., Liu, Y., & Hutter, F. (2021). How Powerful are Performance Predictors in Neural Architecture Search? *Advances in Neural Information Processing Systems*, 34, 28454–28469. <https://arxiv.org/pdf/2104.01177>
- Wistuba, M., & Pedapati, T. (2020). Learning to Rank Learning Curves.
- Wu, J., Cao, M., Shan, L., & Yang, Q. (n.d.). Higher Performance for AutoML: The Benefit of Various Ensemble Bayesian Optimization Strategy. https://github.com/SupUnicorn/bbo_challenge_4th.
- Yan, S., White, C., Savani, Y., & Hutter, F. (2021). NAS-Bench-x11 and the Power of Learning Curves. *Advances in Neural Information Processing Systems*, 27, 22534–22549. <https://arxiv.org/pdf/2111.03602>
- Zheng, X., Ji, R., Chen, Y., Wang, Q., Zhang, B., Chen, J., Ye, Q., Huang, F., & Tian, Y. (2021). MIGO-NAS: Towards Fast and Generalizable Neural Architecture Search. *IEEE Trans-*

actions on Pattern Analysis and Machine Intelligence, 43(9), 2936–2952. <https://doi.org/10.1109/TPAMI.2021.3065138>

Zimmer, L., Lindauer, M., & Hutter, F. (2020). Auto-PyTorch Tabular: Multi-Fidelity MetaLearning for Efficient and Robust AutoDL. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 43(9), 3079–3090. <https://doi.org/10.1109/TPAMI.2021.3067763>

Zoph, B., & Le, Q. V. (2016). Neural Architecture Search with Reinforcement Learning. *5th International Conference on Learning Representations, ICLR 2017 - Conference Track Proceedings*. <https://arxiv.org/abs/1611.01578v2>

8.1 Appendix

8.1.1 Appendix A: Learning Curves vs Forecast

Table 8.1: Condensed EBE-NAS Experiment Results per Dataset

data_id	naml_best_validation_metric	naml_test_accuracy	ebe_decision	ebe_val_acc	ebe_test_acc	ebe_surpassed_val_baseline	ebe_surpassed_test_baseline	ebe_time_used_s
3	0.9941	0.9969	False	0.9941	0.9891	True	False	61.3809
12	0.9850	0.9725	False	0.9750	0.9550	False	False	61.4269
23	0.5619	0.5153	True	0.6017	0.4881	True	False	9.0699
31	0.7575	0.7300	True	0.7562	0.7100	False	False	19.7634
54	0.8515	0.8412	True	0.8824	0.7529	True	False	9.3394
181	0.6050	0.6364	True	0.6345	0.5690	True	False	6.6603
188	0.6661	0.6622	True	0.6780	0.6554	True	False	5.2224
1049	0.9179	0.8836	True	0.9060	0.9110	False	True	19.5543
1067	0.8870	0.8412	True	0.8787	0.8318	False	False	27.6183
1111	0.9822	0.9822	False	0.9822	0.9822	True	True	86.5667
1169	0.6579	0.6607	True	0.6256	0.6293	False	False	93.6522
1457	0.8133	0.8600	True	0.5708	0.4933	False	False	177.8856
1461	0.9133	0.9060	True	0.8831	0.8830	False	False	80.0106
1464	0.7817	0.7733	True	0.8083	0.7867	True	True	3.7188
1468	0.9353	0.9398	True	0.9769	0.9167	True	False	7.0944
1475	0.6047	0.6258	True	0.5444	0.5368	False	False	68.5845
1486	0.9699	0.9697	False	0.9476	0.9504	False	False	90.4251
1487	0.9491	0.9310	True	0.9581	0.9231	True	False	19.1388
1489	0.9210	0.8918	False	0.8960	0.8788	False	False	73.6999

Continued on next page

Table 8.1: Condensed EBE-NAS Experiment Results per Dataset

data_id	naml_best_validation_metric	naml_test_accuracy	ebe_decision	ebe_val_acc	ebe_test_acc	ebe_surpassed_val_baseline	ebe_surpassed_test_baseline	ebe_time_used_s
1494	0.8911	0.8720	True	0.8876	0.8768	False	True	20.5712
1515	0.9087	0.9130	True	0.8261	0.7652	False	False	25.4087
1590	0.8734	0.8774	False	0.8237	0.8310	False	False	90.7140
1596	0.9507	0.9551	False	0.7219	0.7233	False	False	237.2460
4134	0.7939	0.8362	True	0.8050	0.7963	True	False	36.5215
4135	0.9502	0.9469	False	0.9420	0.9422	False	False	73.4512
4534	0.9755	0.9724	False	0.9700	0.9679	False	False	71.5146
4538	0.6709	0.6911	True	0.5652	0.5727	False	False	69.8245
4541	0.6131	0.6152	True	0.5851	0.5825	False	False	90.6034
23512	0.7197	0.7260	True	0.7161	0.7178	False	False	90.3425
23517	0.5233	0.5205	True	0.5059	0.5061	False	False	90.4725
40498	0.6718	0.6827	True	0.5778	0.6153	False	False	62.7789
40668	0.8272	0.8260	False	0.6609	0.6563	False	False	90.3981
40670	0.9582	0.9483	False	0.9529	0.9310	False	False	62.0052
40685	0.9998	0.9995	False	0.7860	0.7859	False	False	88.2175
40701	0.9604	0.9630	False	0.8588	0.8590	False	False	61.7084
40900	0.9967	0.9922	True	0.9951	0.9902	False	False	60.9084
40975	0.9949	0.9971	False	0.9783	0.9682	False	False	64.0157
40978	0.9790	0.9802	False	0.9886	0.9756	True	False	63.4900
40981	0.9081	0.8043	True	0.9189	0.7464	True	False	6.5967
40982	0.7878	0.7738	True	0.7685	0.7147	False	False	54.8814
40983	0.9819	0.9845	True	0.9935	0.9897	True	True	28.4800
40984	0.9395	0.9481	False	0.9216	0.9264	False	False	62.0967
40996	0.8971	0.9024	False	0.7740	0.7631	False	False	227.8870

Continued on next page


Table 8.1: Condensed EBE-NAS Experiment Results per Dataset


data_id	naml_best_validation_metric	naml_test_accuracy	ebe_decision	ebe_val_acc	ebe_test_acc	ebe_surpassed_val_baseline	ebe_surpassed_test_baseline	ebe_time_used_s
41027	0.9007	0.8964	False	0.8190	0.8101	False	False	90.5901
41138	0.9938	0.9934	False	0.9819	0.9819	False	False	90.7054
41142	0.7489	0.7565	True	0.7428	0.7223	False	False	73.4481
41143	0.8187	0.8291	True	0.8096	0.7672	False	False	39.6664
41144	0.8330	0.8551	True	0.6322	0.6178	False	False	61.4547
41145	0.7596	0.7686	True	0.7342	0.6847	False	False	52.0210
41146	0.9524	0.9424	False	0.8951	0.8985	False	False	65.6602
41147	0.6868	0.6853	False	0.5000	0.5000	False	False	143.3326
41150	0.9424	0.9424	False	0.9296	0.9255	False	False	95.5820
41156	0.8444	0.8554	True	0.8464	0.8470	True	False	59.7082
41157	0.8000	0.8000	True	0.8750	0.8000	True	True	3.0432
41158	0.9525	0.9398	False	0.9366	0.9144	False	False	64.2103
41159	0.8228	0.8320	False	0.6012	0.6000	False	False	119.5335
41161	0.9969	0.9972	False	0.9903	0.9922	False	False	144.7661
41162	0.8879	0.8858	True	0.8770	0.8770	False	False	90.5911
41163	0.9842	0.9860	False	0.9419	0.9415	False	False	160.3988
41164	0.7183	0.7130	False	0.6472	0.6292	False	False	65.7742
41165	0.5092	0.5320	True	0.3725	0.3380	False	False	1023.5956
41166	0.6788	0.6911	False	0.5175	0.5111	False	False	90.8297
41167	0.6986	0.5241	False	0.0059	0.0059	False	False	320.6525
41168	0.7105	0.7195	True	0.7082	0.7053	False	False	90.6231
41169	0.3751	0.3723	True	0.0729	0.0713	False	False	93.1162
42732	0.8784	0.8784	False	0.8784	0.8784	True	True	142.6315
42733	0.8360	0.8358	True	0.8317	0.8315	False	False	77.5396


Continued on next page


Table 8.1: Condensed EBE-NAS Experiment Results per Dataset


data_id	naml_best_validation_metric	naml_test_accuracy	ebe_decision	ebe_val_acc	ebe_test_acc	ebe_surpassed_val_baseline	ebe_surpassed_test_baseline	ebe_time_used_s
42734	0.7500	0.7491	True	0.7158	0.7157	False	False	90.7542
42742	0.9636	0.9636	False	0.9636	0.9636	False	True	117.7782
42746	1.0000	1.0000	False	0.5732	0.5732	False	False	207.5776
42769	0.7344	0.7355	False	0.5299	0.5299	False	False	208.1121

























8.1.2 Appendix B: Scatterplot of the Forecast Performance





8.1.3 Appendix C: Error Histogram of the Forecast performance





8.1.4 Appendix D: ShouldYouEvenNN Generational Dynamics

